

Computer Vision hw3

110065502 李杰穎

Problems

1. Finish the rest of the codes for Problem 1 and Problem 2 according to the hint.



```
1 # TODO: Change the output of the model to 10 class.
2 model.fc = nn.Linear(model.fc.in_features, 10)
3 model = model.to(device)
```



```
1 # TODO: Fill in the code cell according to the pytorch tutorial we gave.
2 loss_fn = nn.CrossEntropyLoss()
3 optimizer = optim.Adam(model.parameters(), lr=1e-3)
```



```
1 def train(dataloader, model, loss_fn, optimizer):
2     num_batches = len(dataloader)
3     size = len(dataloader.dataset)
4     epoch_loss = 0
5     correct = 0
6
7     model.train()
8
9     for X, y in tqdm(dataloader):
10         X, y = X.to(device), y.to(device)
11
12         # Compute prediction error
13         pred = model(X)
14         loss = loss_fn(pred, y)
15
16         # Backpropagation
17         optimizer.zero_grad()
18         loss.backward()
19         optimizer.step()
20
21         epoch_loss += loss.item()
22         pred = pred.argmax(dim=1, keepdim=True)
23         correct += pred.eq(y.view_as(pred)).sum().item()
24
25     avg_epoch_loss = epoch_loss / num_batches
26     avg_acc = correct / size
27
28     return avg_epoch_loss, avg_acc
```



```
1 def test(dataloader, model, loss_fn):
2     num_batches = len(dataloader)
3     size = len(dataloader.dataset)
4     epoch_loss = 0
5     correct = 0
6
7     model.eval()
8
9     with torch.no_grad():
10         for X, y in tqdm(dataloader):
11             X, y = X.to(device), y.to(device)
12
13             pred = model(X)
14
15             epoch_loss += loss_fn(pred, y).item()
16             pred = pred.argmax(dim=1, keepdim=True)
17             correct += pred.eq(y.view_as(pred)).sum().item()
18
19     avg_epoch_loss = epoch_loss / num_batches
20     avg_acc = correct / size
21
22     return avg_epoch_loss, avg_acc
```

```

1 epochs = 50
2 train_loss_list = []
3 train_acc_list = []
4 test_loss_list = []
5 test_acc_list = []
6 best_acc = 0
7
8 for epoch in range(epochs):
9     train_loss, train_acc = train(train_dataloader, model, loss_fn, optimizer)
10    # train_loss, train_acc = train(
11    #     sixteenth_train_dataloader, model, loss_fn, optimizer)
12    # train_loss, train_acc = train(
13    #     half_train_dataloader, model, loss_fn, optimizer)
14
15    test_loss, test_acc = test(valid_dataloader, model, loss_fn)
16
17    if test_acc > best_acc:
18        best_acc = test_acc
19        # torch.save(model.state_dict(), 'resnet18_all_IMAGENET.pth')
20        # torch.save(model.state_dict(), 'resnet18_sixteenth_IMAGENET.pth')
21        # torch.save(model.state_dict(), 'resnet18_half_IMAGENET.pth')
22        torch.save(model.state_dict(), 'resnet50_all_IMAGENET.pth')
23        # torch.save(model.state_dict(), 'resnet50_sixteenth_IMAGENET.pth')
24        # torch.save(model.state_dict(), 'resnet50_half_IMAGENET.pth')
25
26    train_loss_list.append(train_loss)
27    train_acc_list.append(train_acc)
28    test_loss_list.append(test_loss)
29    test_acc_list.append(test_acc)
30
31    print(f"Epoch (epoch + 1:2d): Loss = {train_loss:.4f} Acc = {train_acc:.2f} Test_Loss = {test_loss:.4f} Test_Acc = {test_acc:.2f}")
32    print("Done!")
33    print(f"Best Accuracy: {best_acc:.2f}")

```

3. Achieve the best performance given all training data using whatever model and training strategy.

Model: ConvNeXt_Base

Optimizer: AdamW

Scheduler: CosineAnnealingWarmRestarts

Accuracy: 93%

```

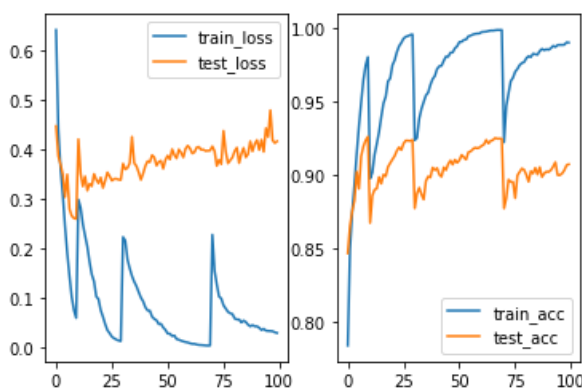
1 import torchvision.models
2 model = models.convnext_base(weights=models.ConvNeXt_Base_Weights.IMAGENET1K_V1)
3 model.classifier[2] = nn.Linear(model.classifier[2].in_features, 10)
4 model = model.to(device)

```

```

1 loss_fn = nn.CrossEntropyLoss()
2 optimizer = optim.AdamW(model.parameters(), lr=1e-3, weight_decay=1e-4)
3 scheduler = optim.lr_scheduler.CosineAnnealingWarmRestarts(optimizer, T_0=10, T_mult=2, eta_min=1e-6)

```

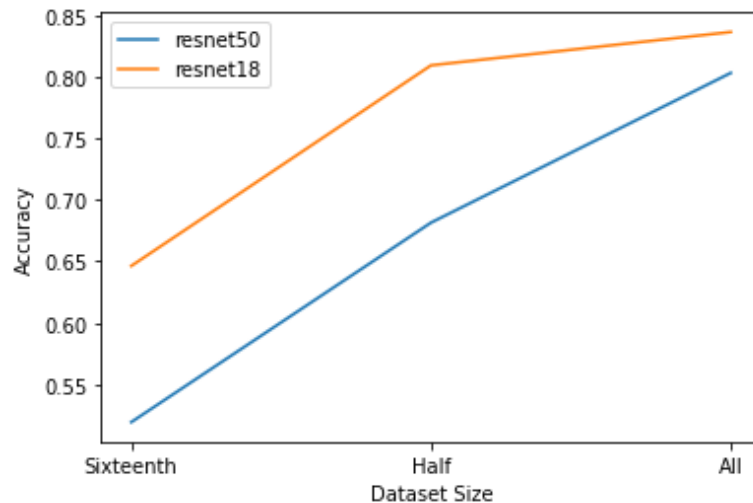


Discussion

- The relationship between the accuracy, model size, and the training dataset size.

The larger the amount of data, the higher the accuracy.

The larger the size of the model, accuracy is not necessarily higher.



- What if we train the ResNet with ImageNet initialized weights (weights="IMAGENET1K V1"), how would the relationship change?

In general, using pretrained weight will achieve better accuracy, and the deeper model will train better than the shallow model.

